

# 1. Problem statement

# 1.1. Understand various dimensions of customer behavior quickly

## Business challenges

- Companies want to understand various dimensions of customer behavior, at “segment-of-one” level. Such dimensions might include: **Frequency, Re-purchase behavior, Future orders**, etc.
- Fast business cycle **requires speedy implementation of ML models**, which can take 2-3 weeks to build from scratch. Slow development process could lead to major business opportunity loss.

## Technical challenges

- With the availability of fast boosted tree algorithms (Catboost, Lightgbm, etc.), feature engineering has become the most important step in the ML modeling process. However, **feature engineering is a manual, often time-consuming process**, which requires a lot of analysis but doesn't always yield the best results.
- Traditional ML models run independently and **cannot learn from each other**, which prevents generalization and induces overfitting. This is potentially problematic when there's only a small amount of labels



### Proposed Problem Statement:

- Building models that predict: **Time to next order (<1week, 1-2weeks, 2-4weeks, >=30days)**, and **Reorder rate of the next order (<10%, 10-50%, 50-90%, >90%)** (= number of re-ordered product / total number of products in the order).
- Using traditional ML models (require feature engineering) as benchmark and experimenting with a sequence-based deep learning model (does not require feature engineering) to speed up development process.
- This project use data from the [Instacart dataset](#).

## 1. 2. Data description



The Instacart datasets contains the following data:

1. Historical orders for each customer, including time gaps between orders
2. Products in each order
3. Metadata of the products (department – product category level 1, aisles – product category level 2)

Data size:

- Number of customers: 206,209
- Number of transactions: 32,434,489
- Range of gaps between orders: 0 days – 30 days

A few summary of the dataset



## 2. Modeling

## 2.1. Feature-based approach: Data processing and modeling

### Data preprocessing methodology

Since I only have access to raw data, I attempt to engineer features with 2 approach:

- Hypothesis based features: Calculate features that make sense business-wise, focus on historical patterns of frequency and re-order behavior (12 features)
  - Build user-product interaction matrix, then apply PCA to extract features with the highest variance. (42 features)
- ⇒ Result in a dataset of 54 features in total

Train-test split: Set aside 10% as test data, apply cross-validation to the 90% training data



Decision made:

- User-product interaction matrix was too large, couldn't fit in memory to perform PCA => Decided to use the user-department interaction matrix and the user-aisle interaction matrix instead.
- Re-formulated the problem of Time to next purchase from regression problem to classification problem due to data distribution (the data was capped at 30 days)

### Modeling methodology

5-fold cross validation to tune hyperparameters for the most popular models:

- Logistic regression
- Random Forest
- Neural network (multi-layer perceptron)
- Boosted tree

Evaluation matrix: Macro F1-score as main matrix, also keep track of ROC curve's AUC



Decision made:

- Used RandomSearch instead of GridSearch because not all hyperparameters are equal
- Dropped a few models due to slow training time or slow inference time: Xgboost, SVM, KNN
- Used LightGBM instead of Xgboost as representation of boosted tree for its significantly faster training time
- Chose Macro F1-score (instead of micro) due to unbalanced data

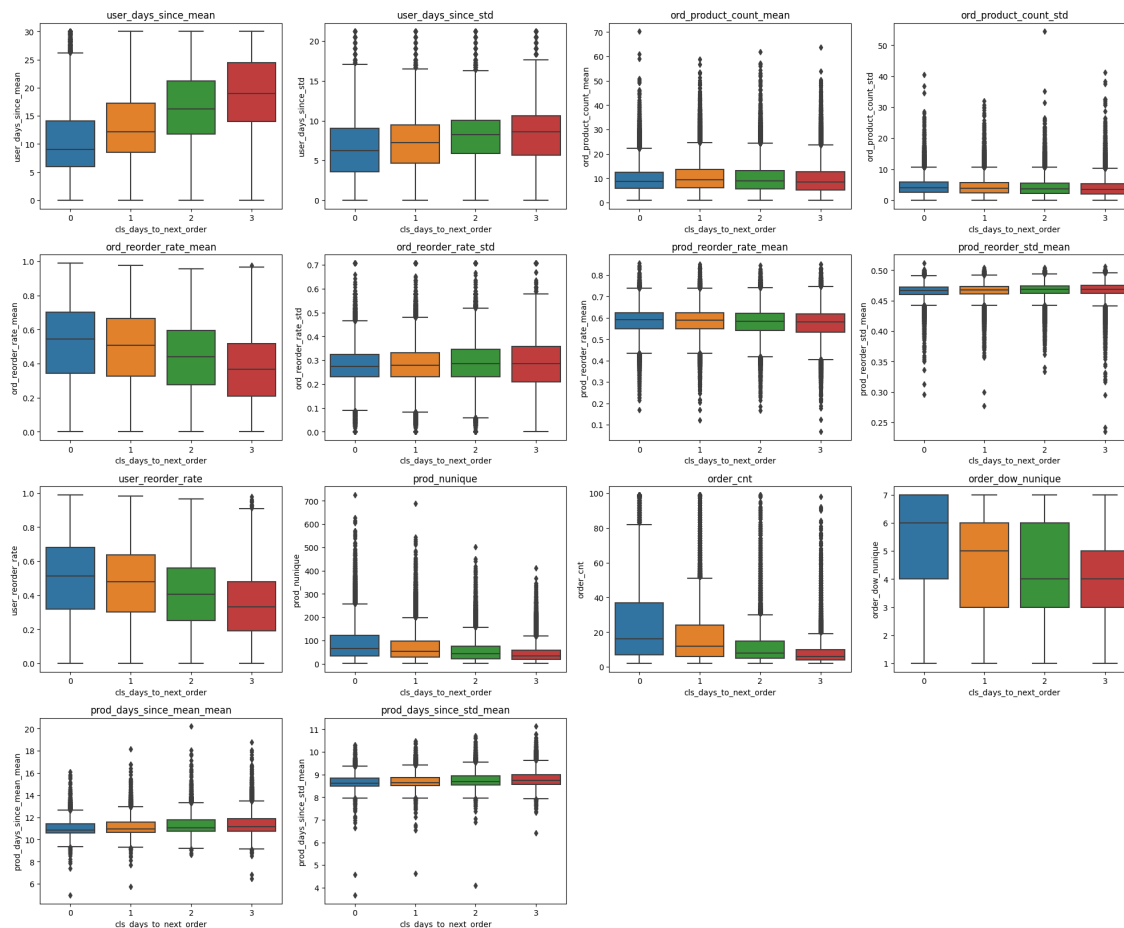
## 2.1. Feature-based approach: Feature engineering results (1/2)

### Hypothesis-based features and correlation with task 1- Time to next order

Overall, we can see that user\_\* features and ord\_\* features offers strong signals.

Meanwhile prod\_\* features don't seem to differentiate much between customers who have short time to next order and customers who have long time to next order.

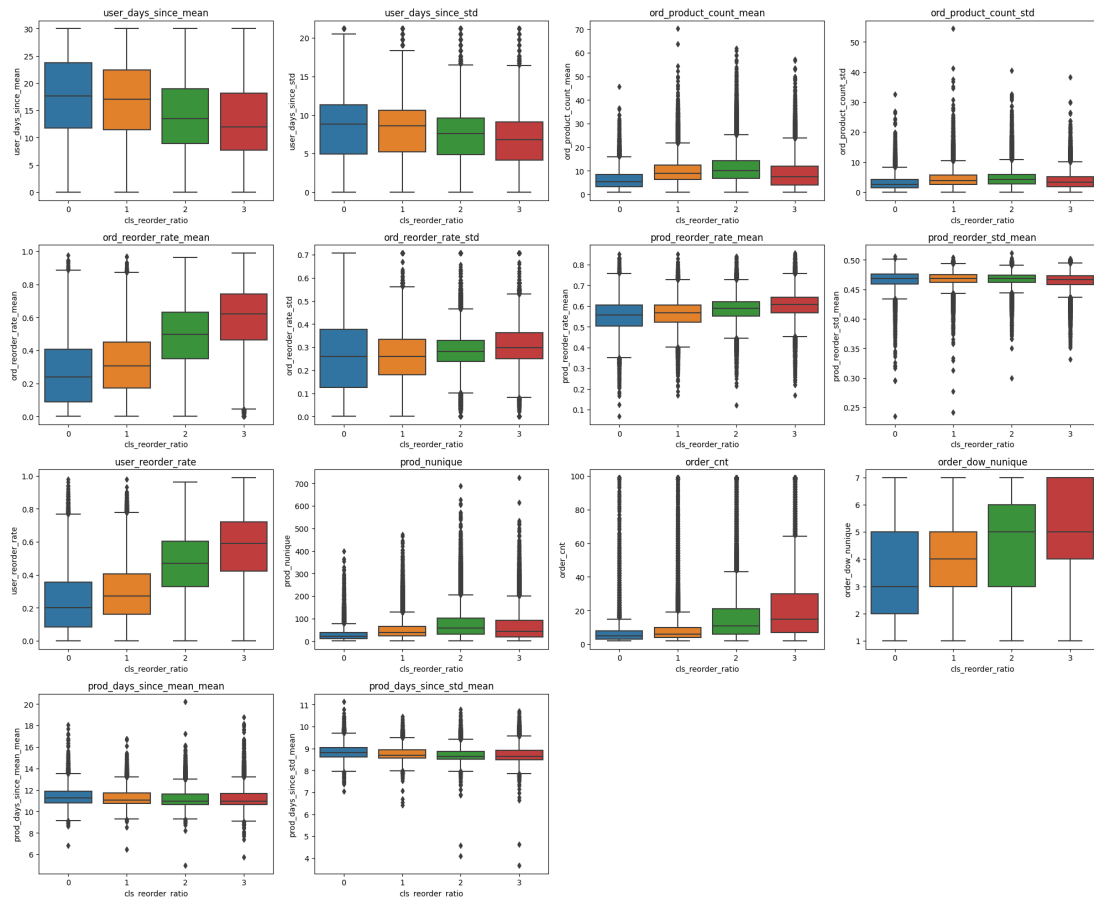
List of 12 features: ['user\_days\_since\_mean', 'user\_days\_since\_std', 'ord\_product\_count\_mean', 'ord\_product\_count\_std', 'ord\_reorder\_rate\_mean', 'ord\_reorder\_rate\_std', 'prod\_reorder\_rate\_mean', 'prod\_reorder\_std\_mean', 'user\_reorder\_rate', 'prod\_nunique', 'order\_cnt', 'order\_dow\_nunique', 'prod\_days\_since\_mean\_mean', 'prod\_days\_since\_std\_mean']



## 2.1. Feature-based approach: Feature engineering results (2/2)

### Hypothesis-based features and correlation with task 2 - Reorder ratio

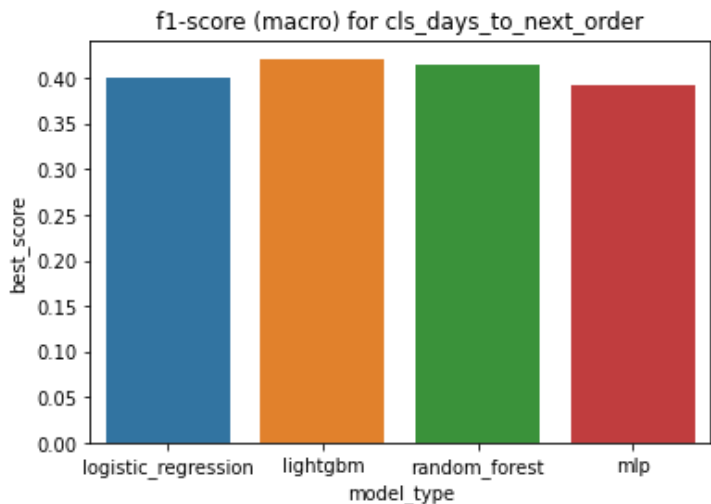
Similar to task 1, we can see that `user_*` features and `ord_*` features seem to differentiate well between customers who have high proportion of reorder products in the next order and customers who have low proportion.



## 2.1. Feature-based approach: Hyperparameter tuning results

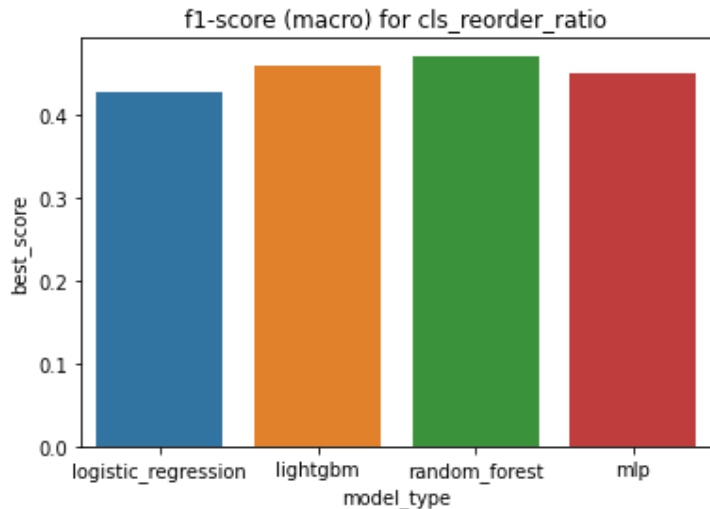
### Task 1 – Time to next order

LightGBM performs the best for this task, achieving a 5-fold cross-validated F1-score of 0.42, while MLP comes last at 0.39



### Task 2 – Reorder rate

RandomForest performs the best for task 2, achieving a cross-validated F1-score of 0.47, while LogisticRegression has the lowest performance at 0.42 macro F1-score





## 2.1. Feature-based approach: Performance on test set – Task 1


Performance metrics on the test set are consistent with cross-validated performance => No overfitting

### LogisticsRegression performance on test set

Evaluating model	precision	recall	f1-score	support
0	0.43	0.55	0.48	4587
1	0.41	0.28	0.33	5762
2	0.29	0.20	0.24	4380
3	0.47	0.62	0.54	5892
accuracy			0.42	20621
macro avg	0.40	0.41	0.40	20621
weighted avg	0.41	0.42	0.40	20621

```
0: 0.7488951630716927
1: 0.6447946226170104
2: 0.609174536849211
3: 0.7490585281946285
```

Some extension of Receiver operating characteristic to multi-class




### LightGBM performance on test set

Evaluating model	precision	recall	f1-score	support
0	0.49	0.51	0.50	4587
1	0.43	0.32	0.36	5762
2	0.30	0.26	0.27	4380
3	0.48	0.62	0.54	5892
accuracy			0.44	20621
macro avg	0.42	0.43	0.42	20621
weighted avg	0.43	0.44	0.43	20621

```
0: 0.7649349013877449
1: 0.6599115686060562
2: 0.6300073746499291
3: 0.7589258795831473
```

Some extension of Receiver operating characteristic to multi-class



## 2.1. Feature-based approach: Performance on test set – Task 1

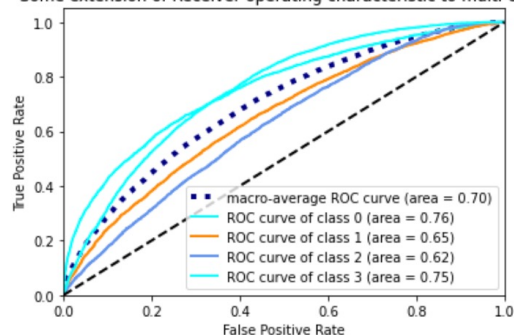
Performance metrics on the test set are consistent with cross-validated performance => No overfitting

### RandomForest performance on test set

Evaluating model	precision	recall	f1-score	support
0	0.50	0.48	0.49	4587
1	0.41	0.33	0.37	5762
2	0.29	0.23	0.25	4380
3	0.46	0.65	0.54	5892
accuracy			0.43	20621
macro avg	0.42	0.42	0.41	20621
weighted avg	0.42	0.43	0.42	20621

0: 0.7573070077622005  
1: 0.6537248819920792  
2: 0.6230309501939817  
3: 0.7530673885201005

Some extension of Receiver operating characteristic to multi-class

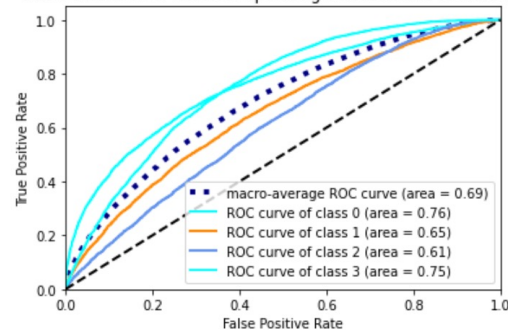


### NeuralNetwork performance on test set

Evaluating model	precision	recall	f1-score	support
0	0.52	0.43	0.47	4587
1	0.41	0.39	0.40	5762
2	0.34	0.07	0.11	4380
3	0.43	0.77	0.56	5892
accuracy			0.44	20621
macro avg	0.43	0.42	0.39	20621
weighted avg	0.43	0.44	0.40	20621

0: 0.7574714310899019  
1: 0.6521413166210603  
2: 0.6108430127370861  
3: 0.7515082515675717

Some extension of Receiver operating characteristic to multi-class



## 2.1. Feature-based approach: Performance on test set – Task 2


Performance metrics on the test set are consistent with cross-validated performance => No overfitting

### LogisticsRegression performance on test set

Evaluating model	precision	recall	f1-score	support
0	0.22	0.61	0.32	1554
1	0.47	0.43	0.45	5380
2	0.63	0.38	0.47	9041
3	0.43	0.55	0.48	4646
accuracy			0.45	20621
macro avg	0.44	0.49	0.43	20621
weighted avg	0.51	0.45	0.46	20621

0: 0.8167574965445632  
1: 0.7441378408709242  
2: 0.6809540265522311  
3: 0.7394854745731769

Some extension of Receiver operating characteristic to multi-class




### LightGBM performance on test set

Evaluating model	precision	recall	f1-score	support
0	0.26	0.54	0.35	1554
1	0.46	0.49	0.47	5380
2	0.62	0.44	0.51	9041
3	0.46	0.53	0.49	4646
accuracy			0.48	20621
macro avg	0.45	0.50	0.46	20621
weighted avg	0.52	0.48	0.49	20621

0: 0.8292921749417266  
1: 0.7456848444166817  
2: 0.6957617944275731  
3: 0.7585339770964237

Some extension of Receiver operating characteristic to multi-class



## 2.1. Feature-based approach: Performance on test set – Task 2


Performance metrics on the test set are consistent with cross-validated performance => No overfitting

### RandomForest performance on test set

Evaluating model	precision	recall	f1-score	support
0	0.34	0.35	0.34	1554
1	0.47	0.58	0.52	5380
2	0.60	0.50	0.55	9041
3	0.48	0.51	0.49	4646
accuracy			0.51	20621
macro avg	0.47	0.48	0.48	20621
weighted avg	0.52	0.51	0.51	20621

0: 0.8253191904264439  
1: 0.7516220554564592  
2: 0.6916736536434768  
3: 0.761338415531694

Some extension of Receiver operating characteristic to multi-class




### NeuralNetwork performance on test set

Evaluating model	precision	recall	f1-score	support
0	0.41	0.21	0.27	1554
1	0.49	0.48	0.49	5380
2	0.56	0.68	0.62	9041
3	0.53	0.41	0.46	4646
accuracy			0.53	20621
macro avg	0.50	0.45	0.46	20621
weighted avg	0.53	0.53	0.52	20621

0: 0.8253651571687971  
1: 0.7554786065467609  
2: 0.6936198156202247  
3: 0.7646930706542792

Some extension of Receiver operating characteristic to multi-class



## 2.2. Sequence-based approach: Foundational concepts

### Representation learning (with neural networks)

#### What it does

Representation learning or feature learning automatically discovers the representations needed for the learning task from raw data.


#### Problem it solves

By eliminating the feature extraction/feature engineering step, representation learning helps:

- Reduce information loss due to feature extraction
- Improve time-to-production of ML models
- Produce good results when domain knowledge is not available

#### Basic principle

Multi-layer neural network learn the representation by learning fundamental concepts in the shallower layers and high-level concepts in deeper layers



### Multi-task learning

#### What it does

Multi-task learning solves multiple tasks at the same time, while exploiting commonalities and differences across tasks.

#### Problem it solves:


- Solve many different problems at the same time
- Reduce or even eliminate overfitting
- Improve the model's generalization a lot by not just considering one particular aspect

#### Basic principle

Theories on why multi-task learning works includes:

- Transfer knowledge between tasks by using some form of shared representation => Improve learning ability
- Introduce a better form of regularization compare to L1, L2 => reduce overfitting

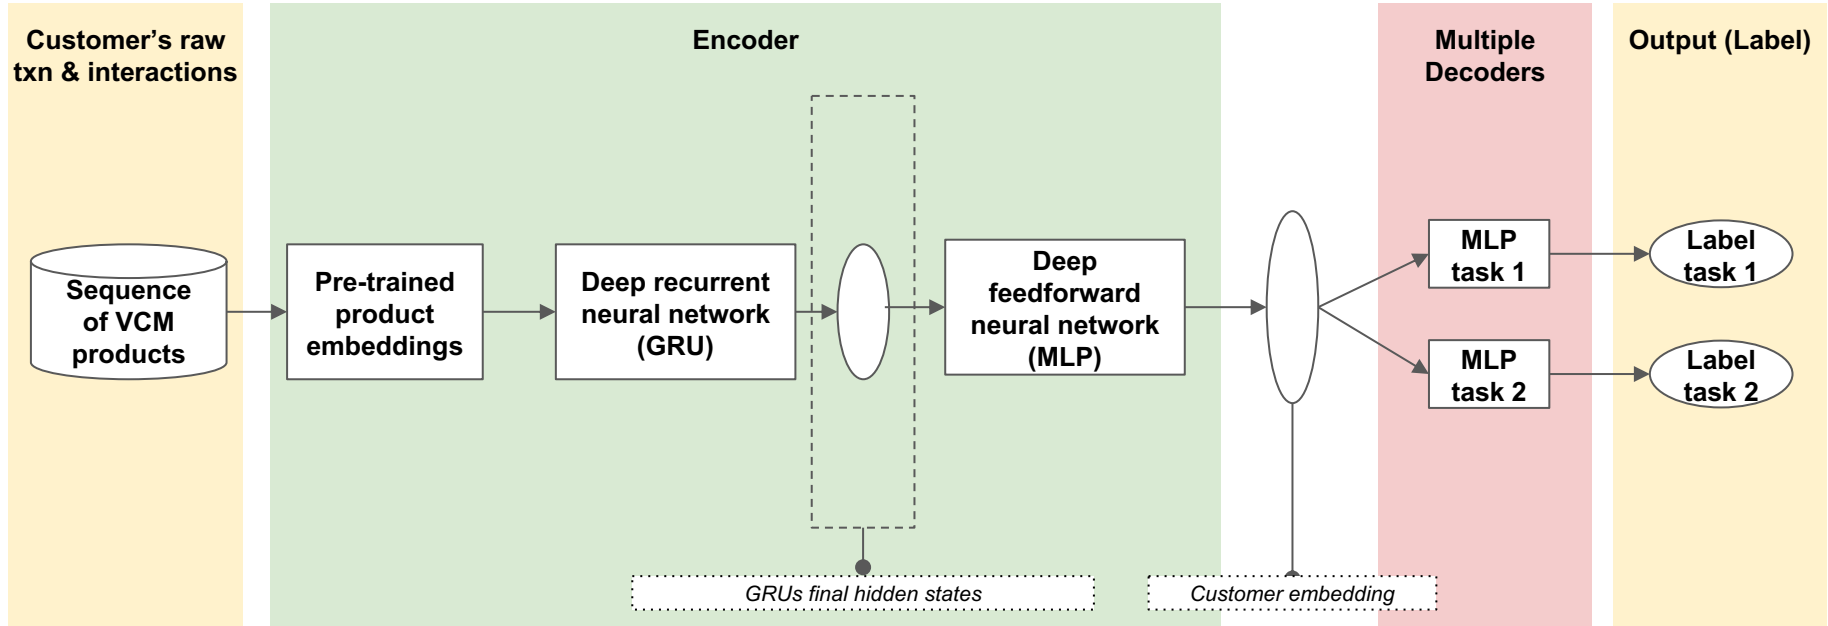
In general, multi-task learning is deemed to resemble human learning better compared to single-task learning



## 2.2. Sequence-based approach: Network architecture

Inspired by Sutskever (2014) which describe the proposed neural machine translation architecture as a model that wants to work, the model uses *Sequence encoding* to encode customer's interactions and *Multilayer perceptron (MLP)* to decode different tasks

**Note:** Due to time constraints, no model tuning was performed for the sequence approach.



## 2.2. Sequence-based approach: Performance on test set

Model's performance for Time to next order

	precision	recall	f1-score	support
0	0.43	0.44	0.43	4484
1	0.27	0.36	0.31	4321
2	0.21	0.28	0.24	3388
3	0.66	0.46	0.54	8428
accuracy			0.40	20621
macro avg	0.39	0.38	0.38	20621
weighted avg	0.45	0.40	0.42	20621

For task 1 – Time to next order, F1-score is comparable with baseline, only slightly lower (0.38 compared to baseline models ranging between 0.39 and 0.42).

Model's performance for Reorder rate

	precision	recall	f1-score	support
0	0.57	0.21	0.31	4185
1	0.41	0.42	0.41	5290
2	0.36	0.58	0.45	5699
3	0.48	0.41	0.44	5447
accuracy			0.42	20621
macro avg	0.46	0.40	0.40	20621
weighted avg	0.45	0.42	0.41	20621

For task 2 – Reorder rate, the gap is more significant (Sequence-based model's F1-score is 0.40 while baseline models range between 0.43 and 0.48).



- Sequence-based models perform slightly worse than well-tuned feature-based models. This is promising for an out-of-the-box deep-learning model.
- With more efforts put into hyperparameter tuning, the performance might be improved to be on-par with feature-based approach.
- Regarding multi-task learning aspect of the model, with only 2 tasks incorporated, multi-task learning's advantages were not expressed significantly.

# 3. Summary and discussion



## 3.1. Summary

### Pros & Cons of the model

Pros and cons of Sequence-based approach compared to feature-based approach:

#### Pros:

- Save weeks spent on feature engineering with built-in representation learning
- Potential to incorporate predictor data from different sources with a plug-and-play mechanism of encoders and decoders
- Produce **product embeddings** and **user embeddings**, which can be re-used easily by other models

#### Cons:

- Is a black-box model, cannot easily be explained like tree-based models or linear models
- Is a big model and therefore take a lot of resources (time, computational resources) to train and tune

### Main contribution of this research

The main contributions of this research are:

1. Combining **representation learning and multi-task learning** into a single model
2. Proposing a **encoder architecture** for some of the most unstructured, hard-to-mine data (products sequence)
3. Demonstrating through experimentation that the model **produces acceptable performance**, not much worse compared to baseline models



## 3.2. Potential future work

### The limits of the current models

- Training time is too slow on a single GPU. Take 30 minutes to train a single model, 50x slower compared to the slowest of feature-based approaches (RandomForest)
- Data input only consider products bought as sequences and does not including the monetary values or product metadata.

### Unexplored territories:

- Currently use Hard Parameter Sharing (HPS) between encoders, there are other multi-task learning architectures to be explored
- Currently use Equal weight to calculate loss function, other weighting strategies that can take into account the importance of each task
- Other training scheduler can make training faster, focusing on rate of improvement for each task
- Task grouping might help improve performance by having different level of parameter sharing

### Future work

1. Reduce training time by experimenting with training schedulers, weighting strategies, and hardware acceleration
2. Improve performance with other multi-task learning architectures
3. Develop encoders to efficiently capture more user information and actions. Some ideas could be:
  - a. Monetary value and product metadata
  - b. Graph embeddings
  - c. Convolutional neural network
4. Apply to real problems



Thank you